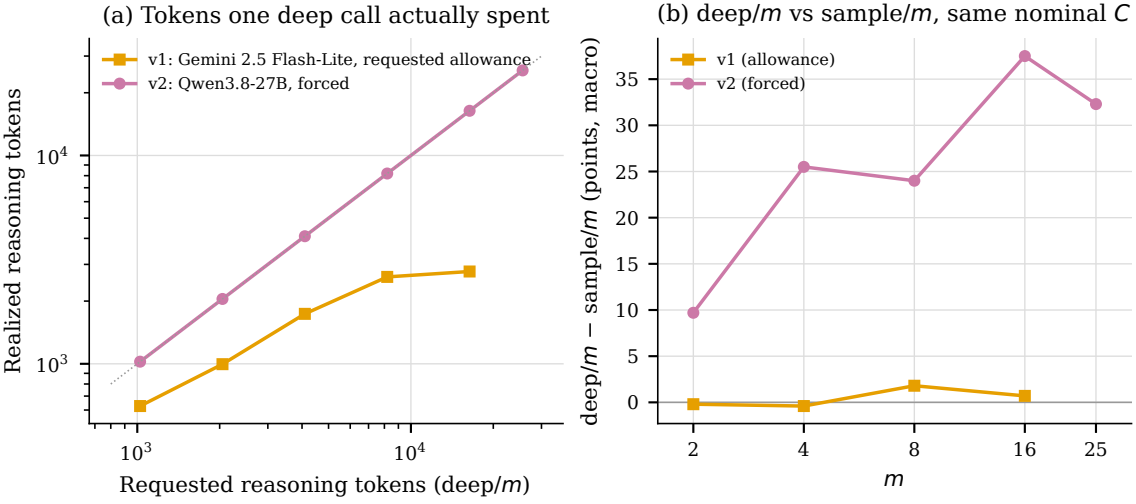
(a) Tokens one deep call actually spent
v1: Gemini 2.5 Flash-Lite, requested allowance
v2: Qwen3.8-27B, forced
Realized reasoning tokens
$10^4$
$10^3$
$10^3$
$10^4$
Requested reasoning tokens (deep/$m$)
(b) deep/$m$ vs sample/$m$, same nominal $C$
v1 (allowance)
v2 (forced)
deep/$m$ − sample/$m$ (points, macro)
35
30
25
20
15
10
5
0
2
4
8
16
25
$m$