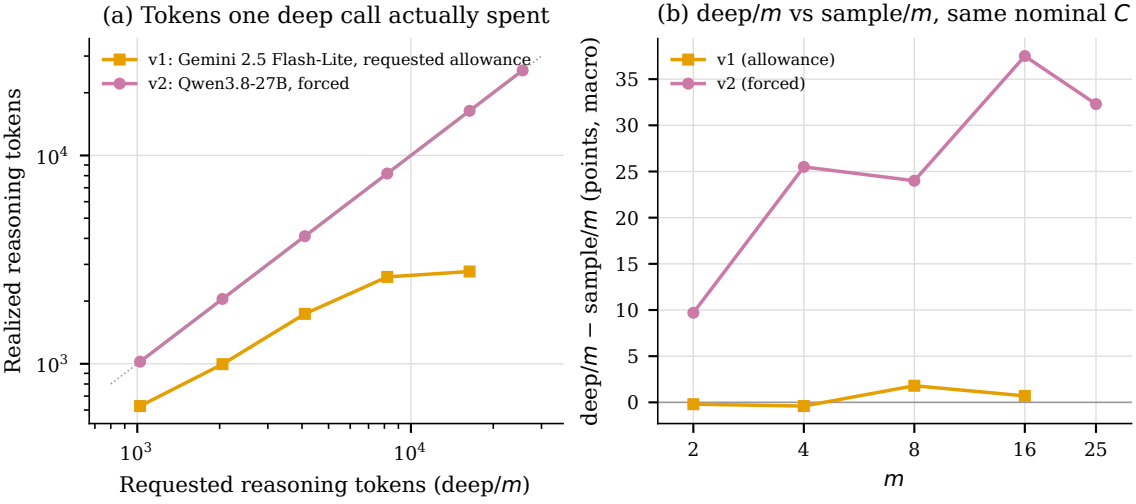
(a) Tokens one deep call actually spent
v1: Gemini 2.5 Flash-Lite, requested allowance
v2: Qwen3.8-27B, forced
Realized reasoning tokens
10^4
10^3
10^3
10^4
Requested reasoning tokens (deep/m)
(b) deep/m vs sample/m, same nominal C
v1 (allowance)
v2 (forced)
deep/m − sample/m (points, macro)
35
30
25
20
15
10
5
0
2
4
8
16
25
m